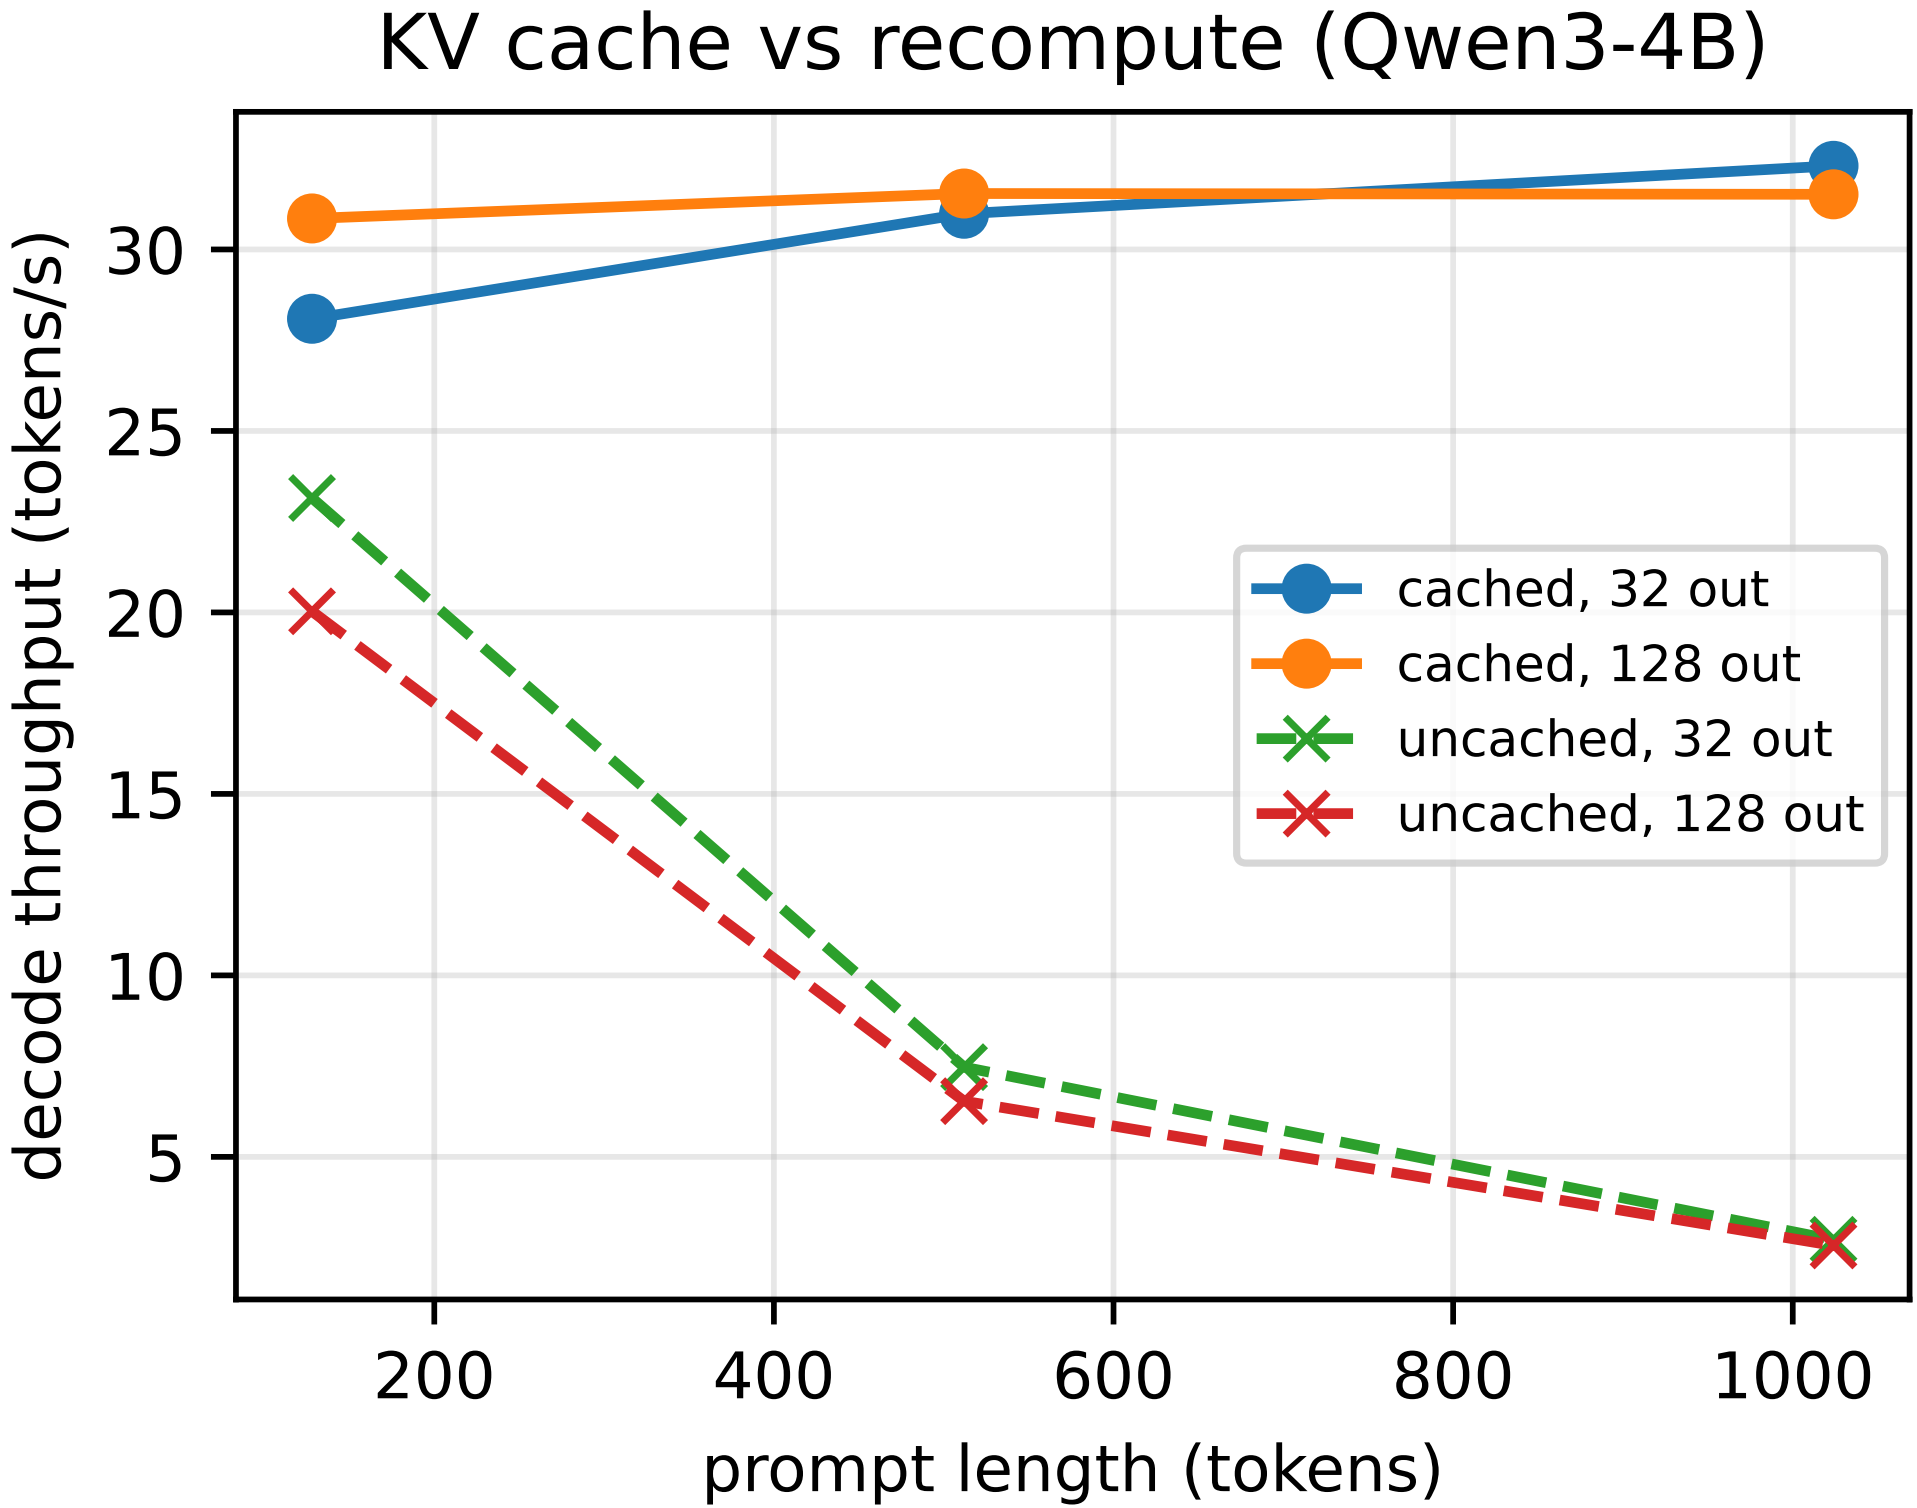
KV cache vs recompute (Qwen3-4B)
decode throughput (tokens/s)
prompt length (tokens)
5
10
15
20
25
30
200
400
600
800
1000
cached, 32 out
cached, 128 out
uncached, 32 out
uncached, 128 out